

NLP Extraction & Tagging Test Pack

Document multilingue (français, anglais, arabe) conçu pour valider : extraction PDF, segmentation layout, détection de langue, tokenisation, POS tagging (NN/VB/JJ...), lemmatisation, NER et détection de bruit.

Mode d'emploi : exécute ton pipeline sur ce PDF. Compare la sortie (sentences_layout / chunks / tokens) avec les attentes ci-dessous. Le header/footer est volontairement répétitif (avec chiffres) pour tester la détection de bruit et le masquage de digits.

- Easy : phrases simples + dates (2026-02-12, 12/02/2026), emails, URLs, nombres, ponctuation standard.
- Medium : listes numérotées, sections (1.1 / a. / (iv)), camelCase, ALLCAPS, mots collés chiffres/lettres, hyphenation sur saut de ligne.
- Complex : micro-table header multi-colonnes, tableau de lignes (facture), colonnes alignées, séparateurs, blocs adresse, multi-colonnes.
- Noise/Edge : répétitions (header/footer), variations de chiffres, longues séquences de ponctuation, mélange scripts (AR/FR/EN).

Quick sanity strings (should be easy to spot)

Emails: support+qa@example.co.uk, facturation@entreprise.dz, test.user+tag@sub.domain.org

URLs: <https://example.com/path?x=1&y=2#frag> — <http://intranet.local/docs/v1/index.html>

IDs: customerID42, invoiceNumber2026A-0007, HTTPStatus200OK,
sha1=da39a3ee5e6b4b0d3255bfef95601890afd80709

English — Complex sentences, contractions, and tricky tokens

Although the committee, which had been debating for months, finally approved the proposal, the CEO insisted that we shouldn't've shipped anything until the QA lead had signed off.

John's colleague said: "If the Joneses' router isn't working, they won't be able to access the U.S. lab's Ph.D. dataset."

Edge cases: 20-year-old, state-of-the-art, non-functional, re-enter, O'Reilly (do not split like a contraction).

Dates & formats: 2026-02-12, 12/02/2026, Feb 12, 2026, 12 Feb 2026. Numbers: 1,234.56 and 1234,56 (comma decimal).

Contact: support+qa@example.co.uk; phone +44 20 7946 0958; ticket #A-00981/2026; ref: INV-EN-0003.

Address block: 221B Baker Street, London NW1 6XE, United Kingdom. Ship-to: 500 Fifth Ave, New York, NY 10110, USA.

Punctuation torture: Wait... what?! This is fine!!! (maybe); brackets (a) [b] {c}; quotes "smart" and 'plain'.

Glued tokens: abc123def, x=1&y;=2, path/to/file_v2.0.tar.gz, user_name@host, 192.168.1.42:8080.

Numbered / alpha / roman list (layout segmentation)

1. First item: Make sure tokenization keeps "shouldn't've" as a chain or decomposes deterministically.

1.1 Sub-item: Multi.level.numbering should start new sections.

a) Alpha item: The label-only line below should be merged with its content.

(iv) Roman item: Check that (iv) is recognized as a section start.

2) Second item: A long hyphenated word split across lines: internationalisation and multi-lingualization.

Français — Sections, élisions, tirets, adresses

Article 1 — Objet : Ce document sert à tester l'extraction et l'étiquetage grammatical. S'il n'est pas robuste, tu verras beaucoup de NN par défaut.

1.1 Détection d'élisions : d'Air France-KLM, l'aéroport, qu'il n'est pas, jusqu'à, lorsqu'il arrive, puisqu'on l'a dit.

a) Tokens piégeux : aujourd'hui (ne pas découper), presqu'île, quelqu'un, e-mail, co-fondateur, va-nu-pieds.

(iv) Ponctuation : « guillemets », points-virgules ; tirets — et parenthèses (test).

Coordinées : facturation@entreprise.dz, tel +213 (0)21 12 34 56, site <https://exemple.dz/portail?ref=AB12>.

Adresse : 12, Rue Didouche Mourad, BP 123, 16000 Alger, Algérie. Livraison : Zone Industrielle Oued Smar — Hangar 7A.

Montants : Total HT = 1 234,56 DZD ; TVA (19%) = 234,57 ; Total TTC = 1 469,13.

Phrase complexe : Si, malgré l'avis du comité, le directeur — qui était pourtant d'accord — change d'idée, alors il faudra re-documenter l'API, re-tester et re-déployer.

Bloc multi-colonnes (clé / valeur) — observe les grands espaces à l'extraction

Client: Société Exemple SARL	Commande: PO-7781
N° Client: CL-00042	Référence: REF/2026/AL-19
Email: facturation@entreprise.dz	Téléphone: +213 21 12 34 56
Date: 15/02/2026	Statut: PAYÉ

Tables & aligned columns — Invoice-style (mixed FR/EN)

Below are tables designed to trigger table detection, micro-table detection, and wrapped description lines. Separators like '----' are included on purpose.

INVOICE	DATE	PO
INV-2026A-0007	2026-02-15	PO-7781
Currency: DZD	VAT: 19%	Terms: Net 30

Qté	Désignation / Description	P. unitaire	TVA	Montant
2	Service de traitement NLP (pack "pro")	45000,00	19%	107100,00
1	Audit + rapport détaillé — includes: tokenisation, POS, lemma, NER , sentiment	6500,00	19%	77350,00
3	Support (email: support+qa@example.co.uk) / hotline +213 2112004,50	14000,00	19%	42840,00
1	Livraison express (221B Baker Street -> Alger) / tracking: 192.168.0.104@08080 0%	8000,00		
	Note: inter-nationalisation + multi-lingualization appear with manual line breaks.			
	Extra: path/to/file_v2.0.tar.gz ; sha1=da39a3ee... ; HTTPStatus200OK.			
<hr/>				
Total HT		180000,00		
TVA (19%)		34200,00		
<hr/>				
Total TTC		214200,00		
<hr/>				

مېقت تامالع + نیوانع + ماقراؤ + یبرع صن - ۋېبرعالا

دانس إلأو، ئەزىزلىلار، ئەغلىلا فاشتكا، یېبرعالا صنلارا جاھتسا رابتخال مۇمۇم فەلمىلا اذە تانايىكلا ىلىع فەرعتلار، يۈچنلار.

جئاتنلار يېف آجىچىن ظحالتىسى، لەمعتى مل اذى؟ حىچىش لەكشىب ئىمزرۇخلا لەمعتى لە: لاؤس مېسىقتلار يېف ئاطخا وأ.

+213 21 12 34 56: فاتاھ، test.user+tag@sub.domain.org ئىنورتكىلى دىربىلە: لاصتا تانايىب

1,234.56: غلبەملىا، INV-AR-0005: ۋەرۇتافلار مقرىر، 15/02/2026: خىراتلار

رئازىجلا، 16000 رئازىجلا، دارمىن شودىد عراش 12: ناونىغىلار

يە امك ظفحۇت نأ بجى (؟،) لىثم ۋېبرعالا زومرلا: ئەظحالىم.

آلودج ھېشىت) ۋاداھىم بەدمۇم ئەلتىك

ۋېمكىلار	فصولا	رەسەل	ۋەبىرىضلار	يەلامىجىلار
2	صوصىنلىكىت ئەمدە	45000.00	19%	107100.00
	ئەپەپ ماقراؤ نو دې فصوللۇ عېبات رەتس			
1	رئازىجلا Baker Street -> عېرسىن نەھش	8000.00	0%	8000.00

Mixed-script stress line

AR/FR/EN in one line: Algiers — Rue Didouche Mourad — customerID42 — دىربىلە: facturation@entreprise.dz — URL: <https://example.com/x?y=2>